

● WEEK 1 · SESSION 2 · COMPANION E-BOOK

UNDERSTAND AI.

*See the Landscape Clearly — the week
you stop guessing about AI and start
seeing how it actually works.*

SESSION 2

The Gen AI landscape — the six kinds,
and how to choose.

FOR EVERYONE

Written for beginners and professionals at
once.

YOU LEAVE WITH

Your AI Landscape Map — the Week 1
deliverable.

WELCOME TO WEEK 1

WEEKLY OBJECTIVE

You had opinions about AI. By the end of this week you have a working model of how it actually functions — technically, economically, and professionally.

Identity shift — AI Curious → AI Aware

THIS BOOK · SESSION 2

SESSION 2

THE GEN AI LANDSCAPE

The six kinds of generative AI, how image models really work, and how to pick the right tool for a job.

HOW EVERY TOPIC IS BUILT

This is your Week 1 companion — read it alongside the two live sessions. It is written for two people at once: the total beginner who has only ever typed into ChatGPT, and the professional who wants the real mechanics. Skim the plain-English parts; stop on the prompts and the build steps.

WHAT	A plain-English definition of the idea.
WHY	Why it matters to you, at work and in the wallet.
WHERE	Where you actually meet it in day-to-day tools.
WHEN	When to think about it — and when to ignore it.
HOW	How it works under the hood, still in plain words.

Every topic also carries: advantages and disadvantages, a real example, three professional prompts at three levels, step-by-step build-up and execution, and the odd Fun Fact.

WHAT'S IN SESSION 2

The three Session 2 topics — the Gen AI landscape and how to choose a tool — then your Week 1 assignment. Jump to whatever you need. (New to the Claude app? The full interface tour is in the Session 1 book.)

2.1	The Six Faces of Gen AI Text, image, video, audio, code, agents — one family, six very different jobs.	p. 4
2.2	How Image AI Works Diffusion: teaching a model to carve a picture out of pure noise.	p. 7
2.3	Choose the Right Tool The AI landscape, and the one skill that outlasts every tool: picking the right one.	p. 10
★	My AI Landscape Report Your Week 1 deliverable	p. 13

21

SESSION 2

THE SIX FACES OF GEN AI

Text, image, video, audio, code, agents — one family, six very different jobs.

TEXT

Claude · GPT · Gemini

IMAGE

Imagen · DALL·E

VIDEO

Runway · Kling

AUDIO

ElevenLabs · Suno

CODE

Cursor · Claude Code

AGENTS

Claude + tools

WHAT

Generative AI comes in six flavours by what it makes: **text** (Claude, GPT, Gemini), **image** (diffusion models), **video** (frame-by-frame generators), **audio** (text-to-speech, music, voice cloning), **code** (execution-aware assistants), and **agents** (AI that uses tools, remembers, and plans). Each works differently under the hood.

WHY

Knowing the six stops you from reaching for the wrong tool. You would not ask a text model to make a video, or a diffusion model to reason about a contract. Matching the modality to the job is half of being an AI generalist.

WHERE

You already meet them daily: a chat assistant (text), Midjourney or Gemini images, Runway or Kling video, ElevenLabs or Suno audio, Cursor or Claude Code (code), and tool-using assistants like Claude with connectors (agents).

WHEN

Pick the modality by the deliverable. Need words or thinking → text. A picture → image. A clip → video. A voice → audio. Working software → code. A repeatable job done for you → an agent.

HOW IT WORKS

Text and code models predict the next token. Image and video models start from noise and refine it into a picture (diffusion). Audio models convert between text and sound waves. Agents wrap a text model in a loop: it decides an action, uses a tool, reads the result, and repeats until the job is done.

ADVANTAGES

- + One skill set (prompting) transfers across all six
- + You can assemble a whole project from parts
- + New modalities keep expanding what one person can make

DISADVANTAGES

- Each has its own tools, quirks and costs
- Quality varies wildly between modalities
- Video and agents are still the least reliable

EXAMPLE

To make one ad: draft the script with a **text** model, generate the visuals with an **image** model, animate them with a **video** model, add a voiceover with an **audio** model — four modalities, one deliverable.



Fun Fact

An “agent” is just a text model in a loop with tools. The intelligence is the same; the difference is it is allowed to act — click, search, send — not just talk.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

List the 6 types of generative AI (text, image, video, audio, code, agents). For each give a one-line "what it makes" and one example tool, in plain words.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

My goal is launching a product. Map which of the 6 Gen AI modalities I'd use at each step of the journey, name one tool per step, and flag the single step where a human must stay in control. Present as an ordered list.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as an AI solutions architect. For my workflow, design a multi-modal pipeline: which modality handles each stage, the exact hand-off between stages, the tool for each, where a human reviews, and the biggest failure risk. Present as a numbered step-flow with a one-line note per step.

BUILD IT UP — STEP BY STEP

- 1 Write down the deliverable you want (a post, a video, an app).
- 2 Break it into parts: words, pictures, sound, motion, logic.
- 3 Label each part with its modality (text / image / video / audio / code / agent).
- 4 Pick one tool per modality — start with the free option.

EXECUTE IT


- 1 Make the text part first; it usually scripts everything else.
- 2 Generate visuals, then motion, then audio — in that order.
- 3 Keep a human check at every hand-off between modalities.
- 4 Assemble the parts in a simple editor (e.g. Canva) at the end.

2.2 SESSION 2

HOW IMAGE AI WORKS

Diffusion: teaching a model to carve a picture out of pure noise.

Start from noise → guided denoising → image



WHAT

Image models like Imagen, DALL·E and Firefly use **diffusion**. In training they take real images, add noise until it is TV static, and learn to reverse the process. To generate, they start from random noise and “denoise” it step by step, guided by your words, until a picture appears.

WHY

This explains why **prompts matter so much** for images, why the same prompt gives different pictures each time (it starts from different noise), and why fine detail like hands and text used to be hard. Understanding diffusion makes you a far better image prompter.

WHERE

Every text-to-image tool is a diffusion model with a friendly front end — Gemini/Imagen, DALL·E in ChatGPT, Adobe Firefly, Ideogram, Midjourney. Video generators extend the same idea across time.

WHEN

Reach for a diffusion tool whenever you need an original picture — a hero image, a mockup, a concept, a thumbnail. Skip it when you need exact text or precise data; those still need care or a specialist model.

HOW IT WORKS

Your words are turned into an embedding (via a system like **CLIP**) that the model uses as a target. Starting from noise, it predicts “what noise to remove” at each step, nudging the image toward your description across ~20-50 steps, all happening in a compressed “latent space” to keep it fast.

ADVANTAGES

- + Creates original images from a sentence
- + Endless variations from one idea
- + Styles, lighting and composition are all promptable

DISADVANTAGES

- Same prompt ≠ same image (randomness by design)
- Struggles with exact text, counting, and anatomy
- Can echo biases or styles from training data

EXAMPLE

Prompt “a red bicycle leaning on a blue wall, golden-hour light” and each run gives a different but on-theme photo — because each starts from fresh noise and denoises toward your words.

**Fun Fact**

Diffusion models learn by destroying pictures. They add noise to millions of images until they are static, then learn to run the film backwards — creation as un-destruction.

THE VISUAL-PROMPT FRAMEWORK — 8 PARTS OF A PROPER IMAGE / VIDEO PROMPT

1 Subject The single main focus — who or what. e.g. a mid-30s founder at a laptop	2 Object Key props / secondary items it interacts with. e.g. a glowing dashboard, coffee cup
3 Environment Setting, background and location. e.g. calm modern home-office at dawn	4 Situation What is happening — the action or moment. e.g. reviewing an AI report, half-smile
5 Reference Style anchor: medium, era, genre, film look. e.g. editorial tech-magazine photo	6 Colour style Palette, temperature, contrast. e.g. deep navy + warm orange, high contrast
7 Motion style Camera + subject movement (write “still” for images). e.g. slow dolly-in, gentle screen glow	8 Geometry style Composition, lens, angle, framing, aspect ratio. e.g. 35mm, eye-level, rule-of-thirds, 3:2

FILLED EXAMPLES — BUILD EVERY VISUAL THE SAME WAY

STILL IMAGE · ART-DIRECTED KEEP THE LABELS, SWAP THE ORANGE Subject: <u>your product</u> held in one hand. Object: soft props that signal <u>use case</u>. Environment: <u>clean studio / real setting</u>. Situation: <u>the key moment of use</u>. Reference: <u>premium brand</u> product photography. Colour style: <u>2-colour palette</u>, high contrast, soft shadow. Motion style: still. Geometry: 50mm macro, top-down, centered, 1:1.	
MOTION · VIDEO KEEP THE LABELS, SWAP THE ORANGE Subject: <u>hero</u>. Object: <u>what they act on</u>. Environment: <u>location + time of day</u>. Situation: <u>the 3-second story beat</u>. Reference: <u>film / ad</u> look. Colour style: <u>graded palette</u>, cinematic contrast. Motion style: slow dolly-in + subtle parallax, 24fps. Geometry: anamorphic 2.39:1, low angle, leading lines.	

BUILD IT UP — STEP BY STEP

1 Pick your idea, then fill the 8 parts above — one short phrase each.
2 Write “still” under Motion style for an image; describe the move for video.
3 Paste the filled prompt into your image tool and generate 3-4.
4 Keep the parts that worked; change ONE part at a time to steer the rest.

EXECUTE IT

1 Reuse the 8-part skeleton for every visual — it makes results repeatable.
2 For a matching set, keep Subject + Colour style fixed; vary Geometry only.
3 For exact words on the image, put them in “quotes” inside Object.
4 Finish the winner in Canva — add your text and logo there.

23

SESSION 2

CHOOSE THE RIGHT TOOL

The AI landscape, and the one skill that outlasts every tool: picking the right one.

THE JOB	CATEGORY	START WITH
Think, write, orchestrate	Text / hub	Claude
Research with sources	Answer engine	Perplexity
Summarise your own docs	Knowledge base	NotebookLM
Explore open models	Model hub	Hugging Face
Make an image	Image (diffusion)	Gemini / Imagen
Remember across tools	External memory	Notion (via Claude)

Tools change monthly; the **map** does not. Name the job → the category → today's best tool.

WHAT

There are thousands of AI tools, but they cluster into a handful of jobs. The generalist's skill is not knowing every tool — it is a **mental map**: for research use Perplexity, for knowledge synthesis NotebookLM, to explore open models Hugging Face, and Claude as the orchestrating brain that ties them together.

WHY

Tools change monthly; the map does not. If you can name the job and the category, you can always find today's best tool for it. This is exactly what the Week 1 assignment builds — your personal AI Landscape Map.

WHERE

Your map lives wherever you work — a Claude Project, a Notion page, a doc. The tools live across the web; the map is what makes them usable to *you*.

WHEN

Build the map once this week, then update it quarterly. Consult it whenever a new task appears or a shiny new tool tempts you — does it fill a real gap on your map, or just duplicate one you have?

HOW IT WORKS

Start from your role and your recurring tasks, not from the tools. For each task, name the modality (from topic 2.1), then the category (research, writing, image, automation...), then pick one tool. Claude is the hub: it can research, draft, and route work to the others through connectors and skills.

ADVANTAGES

- + A map survives the endless churn of new tools
- + Stops shiny-tool overwhelm and wasted subscriptions
- + Makes you fast: you always know where to start

DISADVANTAGES

- Needs occasional upkeep as tools evolve
- Too many tools still fragments your workflow
- A map is only as good as the tasks you put in it

EXAMPLE

A marketer's map: research → Perplexity; strategy & writing → Claude; images → Gemini; short video → Runway; automation → Zapier; memory → Notion (via Claude connector). Six jobs, six tools, one Claude hub.

**Fun Fact**

The tool you master beats the tool that is technically best. A generalist who knows one hub tool deeply out-ships a dabbler with twenty tabs open.

◆ FIND IT IN CLAUDE

Customize → Connectors

Wire the “right tool” straight into Claude. Once you have named the job and the tool, many of those tools plug in as **connectors** (Drive, Notion, Slack...). Connect them under Customize → Connectors, then reach them from the “+” in any chat.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Ask me 3 quick questions about my work, then recommend the 5 AI tools I should start with and one line on what each is for. Keep it beginner-friendly.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Help me build an AI tool map for my role. List my 8 most common tasks, the AI category for each (research, writing, image, automation...), and one recommended tool per task. Output as a 3-column table.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as my AI strategist. Produce my "AI Landscape Report": 12 tools across 6 categories for my field. For each: what it does, how it works at a high level, who it's best for, how Claude compares, and one surprising use case. Format as a shareable brief with a summary table at the top.

BUILD IT UP — STEP BY STEP

- 1 Open a new Claude Project and name it "My AI Landscape".
- 2 List your role and your 8 most common tasks.
- 3 For each task, write the modality and the category.
- 4 Ask Claude (with Perplexity for facts) to suggest one tool per task.

EXECUTE IT

- 1 Group the tools into 6 categories; drop any duplicates.
- 2 Mark a priority shortlist — the 3-4 you will actually use first.
- 3 Add one "surprising use case" per category.
- 4 Export it as a PDF or Notion page — this is your Week 1 assignment.



YOUR WEEK 1 DELIVERABLE

MY AI LANDSCAPE REPORT

The one thing to build this week — it feeds straight into Week 2.

IMAGE PLACEHOLDER — GENERATE & DROP IN

A hero image for your report cover

STRUCTURED PROMPT FOR CHATGPT / GEMINI / FIREFLY

SUBJECT	a central glowing node labelled “Claude”	OBJECT	six smaller labelled nodes — research, writing, image, video, automation, memory — joined by thin lines
ENVIRONMENT	deep-navy studio backdrop, soft vignette, empty space along the top for a title	SITUATION	a hub-and-spoke map, quietly “powered on”, lines gently glowing
REFERENCE	clean flat-vector infographic, modern tech-brand illustration	COLOUR STYLE	deep navy base, warm orange accents, small teal & purple highlights, high contrast
MOTION STYLE	still (for a video cover: slow parallax drift across the nodes)	GEOMETRY STYLE	centered hub-and-spoke, generous negative space at top, minimal, 3:2

BRIEF	Research and document 12 AI tools across 6 categories relevant to your field. For each: what it does, how it works at a high level, who it is best for, and how Claude compares. Include one “surprising use case” per category.
SUBMIT AS	A structured PDF or Notion document — shareable.
TIME	60–75 minutes
TOOLS	Claude · Perplexity · NotebookLM · Hugging Face
FEEDS INTO	Feeds into Week 2 — you will know exactly which tools to connect into your Claude workspace.

◆ FIND IT IN CLAUDE

Sidebar → Projects → New project

Build the report inside a Project. Create a Project for this report so your chats, sources, and drafts live in one place you can reopen next week. Add Perplexity and Notion as **connectors** (Customize → Connectors) to pull research in without copy-paste.

WEEK 1 OUTCOME

THE CONFUSION IS GONE.

A technical and practical understanding of how AI models work, what the landscape looks like, and where Claude fits.

The confusion is gone. The foundation is real.

AI Curious → AI Aware

NEXT · WEEK 2

You now know **what** the tools are. Week 2 turns that map into a working setup — connecting the right tools into your **Claude workspace** so they actually do the job for you.