

● WEEK 3 · COMPANION E-BOOK

THINK + CREATE.

Prompt Engineering & the AI Creative Studio — the week you stop asking Claude questions and start directing it.

TWO SESSIONS

A prompt-engineering masterclass, then the AI creative studio.

FOR EVERYONE

Written for beginners and professionals at once.

YOU LEAVE WITH

A prompt library and a portfolio-grade AI Media Kit.

WELCOME TO WEEK 3

WEEKLY OBJECTIVE

You were asking Claude questions. This week you engineer precise instructions that produce expert-level output — and direct full visual, audio and motion production from a brief.

Identity shift — Claude Native → AI Creative Director

YOUR TWO SESSIONS

SESSION 1

THE PROMPT ENGINEERING MASTERCLASS

How Claude reads a prompt, the six-part anatomy, and the reasoning, output-control and pipeline frameworks that produce expert output.

SESSION 2

THE AI CREATIVE STUDIO

How image, video and audio models actually work — and how to direct them, prompt by prompt, into one cohesive AI Media Kit.

HOW EVERY TOPIC IS BUILT

This is your Week 3 companion — read it alongside the two live sessions. It is written for two people at once: last week’s operator who configured a Claude workspace, and the professional ready to direct. Skim the plain-English parts; stop on the frameworks, the prompts, and the build steps.

WHAT

A plain-English definition of the idea.

WHY

Why it matters as you become a director, not just a user.

WHERE

Where you actually meet it — in Claude or a creative tool.

WHEN

When to reach for it — and when to ignore it.

HOW

How it works under the hood, still in plain words.

Every topic also carries: advantages and disadvantages, a real example, three professional prompts (or the 8-part visual framework), step-by-step build-up and execution, and the odd Fun Fact.

WHAT'S IN THIS WEEK


A full tour of the Claude app (3 pages), then the eight topics across the two sessions, then your assignment. Jump to whatever you need.

◆	Getting around Claude The whole interface + how to connect & automate	p. 4
1.1	Why Prompting Works How Claude reads a prompt — and the six-part anatomy that makes output expert-level.	p. 7
1.2	The Reasoning Frameworks Zero-shot to Tree-of-Thought — the frameworks that change how Claude thinks.	p. 10
1.3	Controlling the Output Shape, constrain and refine — make Claude produce exactly the form you need.	p. 13
1.4	Prompting as a Pipeline Chain prompts, ground them in sources, and diagnose failure — prompting at system scale.	p. 16
2.1	Image AI How diffusion really works — and the prompt architecture that directs it.	p. 19
2.2	Video AI Frame synthesis, motion and temporal consistency — directing AI video from a scene brief.	p. 22
2.3	Audio & Voice AI Synthetic voice, cloning and music — how it works, and how to direct it responsibly.	p. 25
2.4	3D & Media Kit Assembly Neural 3D, Blender by text command, and assembling every asset into one kit.	p. 28
★	AI Productivity Playbook + Media Kit Your Week 1 deliverable	p. 31

BEFORE YOU START · THE CLAUDE-FIRST TOUR

GETTING AROUND CLAUDE

This whole program runs on Claude. Learn the app once — here is the entire window, every part labelled.



THE CLAUDE APP — LEFT SIDEBAR, TOP BAR, CHAT (CENTRE), MESSAGE BOX (BOTTOM), AND THE LIVE PANEL (RIGHT)

THE LEFT SIDEBAR — WHERE YOU MOVE BETWEEN KINDS OF WORK

New — Start a fresh chat or Cowork session. Nothing carries over unless the session shares a Project.

Artifacts — Your gallery of built pages and docs (dashboards, mini-apps, reports) that live on beyond the chat and can be re-opened or shared.

Dispatch — Hand a job to a background agent that plans and runs it while you do other things. (Beta.)

Projects list — Your projects pinned for quick access; click one to jump straight into its space.

Projects — Your Projects — shared folders of chats, files and instructions that persist across sessions and your team.

Scheduled — Tasks that run on a schedule — hourly, daily, weekly. View, pause, edit or run them here.

Customize — Your home for connectors, custom instructions, styles and preferences — set once, applies everywhere.

Account · plan — Bottom-left: who you are and your plan (e.g. Max). The download icon gets the desktop app.

THE TOP BAR — WHERE YOU ARE, AND HOW TO GET AROUND

Chat · Cowork / Code — The mode switch. “Chat & Cowork” for everyday work with files and tools; “Code” opens Claude Code, the developer surface.

Back / forward — Move through screens you have visited, like a browser.

Panel & theme — Top-right icons: show or hide the side panels and switch light / dark.

Search — Find any past chat, project or output by keyword.

Breadcrumb — Shows where you are (e.g. Playbooks / Frontend design...). Click a crumb to go up a level.

THE MESSAGE BOX — HOW YOU BRIEF CLAUDE AND CONTROL IT

The “+” — Add connectors, upload files, and switch on skills for this message.

Approval: Manual / Auto / Skip — When Claude asks before acting. Manual = approve every action; Auto = auto-approves reads, checks writes; Skip = no prompts (trusted work only).

Model · reasoning — Pick the model (e.g. Opus 4.8) and the thinking effort — High for hard problems, lower for speed.

Add folder — Give the session a folder from your computer to read and write — pick the smallest folder that holds what is needed.

Dictate — The mic — talk instead of type.

THE PANEL ON THE RIGHT — WHAT CLAUDE SHOWS YOU WHILE IT WORKS

Progress — The live task list — every step Claude is taking, ticking off as it goes.

Browser · Claude in Chrome — The browser Claude is driving for web tasks, when a browser is connected.

Connectors — The apps switched on for this session (e.g. Web search); add more from “+” or Customize.

Outputs — The finished files from this session — open, preview, download, or send some to Drive.

Context — What this session can see — the connected Project and folders (e.g. Playbooks).

Skills — The packaged know-how Claude can apply here on demand.

◆ FIND IT IN CLAUDE

Message box → Manual / Auto / Skip

Set the approval mode before a big task. New to Cowork? Start on **Manual** so you approve each action. As you trust a routine, move to **Auto**. Keep **Skip** for small, contained jobs only.

THE THREE THINGS THAT TURN CLAUDE FROM A CHATBOT INTO A WORKSPACE

Anyone can type into Claude. The step up is **connecting your own apps**, **putting work on repeat**, and **handing off jobs to the background**. Here is where each lives and how to switch it on.

1

CONNECT YOUR APPS

CONNECTORS

A connector plugs one of your apps — Gmail, Google Drive, Notion, Slack, and hundreds more — straight into Claude, so it can read and act on that data instead of you copy-pasting.

- 1 Open Settings → Connectors (or, in a chat, click “+” → Connectors → Manage connectors).
- 2 Click the “+” next to Connectors to open the directory, then pick the app you want.
- 3 Click Connect, then sign in and approve access in the pop-up (this is OAuth — you approve, Claude never sees your password).
- 4 Back in any chat, click “+” → Connectors and toggle the app on for that conversation.

Free, Pro, Max, Team & Enterprise. Team/Enterprise: an Owner enables a connector for the org first. Choose the least access that does the job.

2

PUT WORK ON REPEAT

SCHEDULED TASKS

A scheduled task is a prompt Claude runs on its own on a set cadence — a Monday digest, a daily inbox triage — and reports back with no nudge from you.

- 1 In the left sidebar, click Scheduled.
- 2 Click New task → Create with Claude (guided) or Set up manually.
- 3 Write the prompt as a full standalone instruction, pick the cadence (hourly / daily / weekly / weekdays), and set the approval mode + folder.
- 4 Click Schedule. Pause, edit, or run it on demand any time from the same page.

Pro, Max, Team & Enterprise. Desktop for all paid plans; web & mobile in beta.

3

HAND OFF, WALK AWAY

DISPATCH

Dispatch is a background agent. You describe an outcome once; it breaks the job into tasks, runs each on its own, and surfaces the results later — even from your phone.

- 1 In the left sidebar, select Dispatch.
- 2 Describe the outcome like you would brief a colleague, in one message.
- 3 Let it plan — it starts child tasks that each show a status (Running, Awaiting input, Completed...).
- 4 Check back later: open any task to read the transcript and grab the files it produced.

Pro & Max, on the latest Claude Desktop (macOS/Windows). Keep the desktop app open to assign from mobile.

1.1

SESSION 1

WHY PROMPTING WORKS

How Claude reads a prompt — and the six-part anatomy that makes output expert-level.

The six-part prompt anatomy — stack all six for expert output

ROLE	who Claude acts as
CONTEXT	the background it assumes
TASK	the exact ask
CONSTRAINTS	limits & format locks
OUTPUT FORMAT	the shape you want back
VERIFICATION	catch its own errors

WHAT

A prompt is not a search box. Claude turns your words into **tokens**, weighs them with **attention**, and **generates** the most likely useful continuation. A complete prompt has six parts: **Role · Context · Task · Constraints · Output format · Verification**.

WHY

The structure of your input changes the quality of the output more than any tool switch. The same question, asked with the six parts in place, is the difference between a generic answer and an expert one.

WHERE

Every message box — but you feel it most inside a **Project**, where the Role and Context live in the system prompt and each message adds only the Task.

WHEN

Use the full anatomy for anything that matters. For a throwaway question a plain sentence is fine — the skill is knowing which is which.

HOW IT WORKS

Tokenisation splits your text into pieces; attention lets Claude weigh which pieces matter for what comes next; generation predicts forward one token at a time. Give it a clear role and constraints and you are steering all three — narrowing the space of likely outputs toward the one you want.

ADVANTAGES

- + Turns vague answers into expert-level output
- + Repeatable — the anatomy works on any task
- + Makes Claude's behaviour predictable and testable

DISADVANTAGES

- Over-structuring a trivial ask wastes time
- A missing Verification step lets errors through
- Bad Context poisons everything downstream

EXAMPLE

Instead of “write a cold email,” you give Role (“act as a B2B copywriter”), Context (“we sell to CFOs”), Task, Constraints (“under 90 words, no hype”), Output format (“subject + body”) and Verification (“flag any claim you cannot support”). The reply is usable on the first try.

**Fun Fact**

Adding two words — “think step by step” — measurably improves Claude's answers on hard problems. It is not magic: you are telling the generation process to spend tokens reasoning before it commits to an answer.

◆ FIND IT IN CLAUDE[Customize → Skills](#)

Store the six-part prompt as a Skill. Save your best six-part anatomy prompt as a **Skill** so the Role and Output format are always applied. Attach it to a Project and call it with /.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Rewrite my request using the six-part prompt anatomy — Role, Context, Task, Constraints, Output format, Verification. My request is: paste it.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Act as a prompt engineer. Take my task describe and build a full prompt with all six parts. Make the Constraints testable and add a Verification step that catches your own mistakes.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as my prompt architect. For recurring task, design a reusable six-part prompt I can store as a Skill: fixed Role and Output format, slots for Context and Task, and a Verification step tuned to the errors this task usually produces.

BUILD IT UP — STEP BY STEP

- 1 Take a task you do often and write it as one plain sentence.
- 2 Add the six parts: Role, Context, Task, Constraints, Output format, Verification.
- 3 Run both versions and compare the output side by side.
- 4 Keep the six-part version as a template in your Project.

EXECUTE IT

- 1 Make every Constraint concrete and checkable — “under 90 words”, not “short”.
- 2 Always include a Verification line — it is the cheapest quality gain there is.
- 3 Store the Role and Context in the Project system prompt so each message is just the Task.
- 4 Save your best anatomy prompts as a Skill for one-line reuse.

1.2

SESSION 1

THE REASONING FRAMEWORKS

Zero-shot to Tree-of-Thought — the frameworks that change how Claude thinks.

Pick the framework by the job

ZERO-SHOT

just ask

FEW-SHOT

2–5 examples

CHAIN-OF-THOUGHT

reason, then answer

TREE-OF-THOUGHT

weigh 3 routes

WHAT

A set of named techniques for shaping Claude's reasoning: **zero-shot** (no examples), **few-shot** (2–5 examples), **chain-of-thought** (reason before answering), **tree-of-thought** (weigh parallel approaches), **role prompting** (expert personas), and **meta-prompting** (Claude writes and critiques your prompts).

WHY

Each framework unlocks a different kind of quality. Few-shot locks a format; chain-of-thought fixes reasoning errors; role prompting deepens expertise. Knowing which to reach for is the core of the craft.

WHERE

All inside the normal chat — they are ways of writing the prompt, not features to switch on. Store the ones you reuse as Skills.

WHEN

Zero-shot for simple asks; few-shot when format matters; chain-of-thought for logic and math; tree-of-thought for decisions with trade-offs; meta-prompting when you are stuck on how to ask.

HOW IT WORKS

Few-shot works because Claude learns the pattern from your examples in-context — quality and diversity of examples matter more than quantity. Chain-of-thought works because reasoning tokens give the model room to work before answering. Tree-of-thought has it lay out several routes and self-evaluate before committing.

ADVANTAGES

- + A framework for every kind of task
- + Few-shot examples lock format without long instructions
- + Meta-prompting turns Claude into your prompt coach

DISADVANTAGES

- Few-shot with bad examples teaches bad habits
- Chain-of-thought adds length — skip it for simple asks
- More frameworks stacked is not always better

EXAMPLE

A pricing decision: you ask Claude to lay out three approaches (penetration, premium, tiered), reason through each against your constraints, then recommend one with its risk. That is tree-of-thought — and it beats a single confident answer every time.

**Fun Fact**

Few-shot prompting is really “in-context learning” — the model is not retrained, it just infers your pattern from the examples in the window. Three good examples can outperform a paragraph of instructions.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Answer this, but first think step by step and show your reasoning before the final answer: paste the problem.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Here are 3 examples of the output I want paste examples.
Following the same pattern exactly, produce the output for this new input: input. Match the style and structure of the examples.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as a strategist. Use tree-of-thought: lay out 3 distinct approaches to decision, reason through each against these constraints list, score them, and recommend one with the single biggest risk. Then meta-prompt me: what did I leave out?

BUILD IT UP — STEP BY STEP

- 1 Pick one real task and solve it zero-shot first — note where it falls short.
- 2 Add 2–3 strong examples (few-shot) and rerun; watch the format lock in.
- 3 For a reasoning task, add “think step by step” and compare.
- 4 For a decision, ask for 3 approaches weighed against your constraints.

EXECUTE IT

- 1 Choose the framework by the job: format → few-shot, logic → chain-of-thought, trade-offs → tree-of-thought.
- 2 Curate few-shot examples like gold — they teach the pattern.
- 3 When you cannot phrase a prompt, meta-prompt: ask Claude to write it.
- 4 Save each framework as a titled Skill in your prompt library.

1.3

SESSION 1

CONTROLLING THE OUTPUT

Shape, constrain and refine — make Claude produce exactly the form you need.

Control the shape — output you can pipe straight onward

JSON

```
{ task, owner,  
  due, priority }  
  
for tools
```

TABLE

for people

CONSTRAINTS

< 90 words
no hype · CFO
define "good"

WHAT

The techniques that control the *shape* of the answer: **structured output** (JSON, XML, tables, markdown templates), **constraint prompting** (word limits, format locks, negative constraints, audience), **iterative refinement** (critique → revise → validate), and **multi-turn design** (building context across turns).

WHY

A right answer in the wrong shape is still rework. Controlling output turns Claude into a reliable component — output you can paste into a sheet, a form, or another tool without cleanup.

WHERE

In the prompt itself (“return valid JSON with these keys”), and across a conversation you design deliberately rather than stumble through.

WHEN

Use structured output whenever something downstream reads the result; constraints whenever “good” has a definition; refinement loops for anything you would otherwise redraft by hand.

HOW IT WORKS

You tell Claude the exact schema and it fills it; negative constraints (“never use the word...”) remove whole classes of bad output; a refinement loop asks Claude to critique its own draft against your criteria, then revise. Multi-turn design means each message deliberately adds context the next one builds on.

ADVANTAGES

- + Output you can pipe straight into other tools
- + Constraints define “good” so you stop redrafting
- + Refinement loops raise quality without you rewriting

DISADVANTAGES

- Over-tight formats can cut off good reasoning
- A broken schema request yields broken JSON
- Long multi-turn threads drift — reset when they do

EXAMPLE

You ask for “valid JSON: {task, owner, due, priority} for each action item” from a meeting transcript. The result drops straight into your tracker via a Zap — no cleanup, no copy-paste.

**Fun Fact**

Ask for JSON and Claude will usually give you valid JSON — but add “JSON only, no prose” and it stops wrapping it in a friendly sentence that breaks your parser. The constraint is doing real engineering work.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Give me the answer as a markdown table with these columns:
columns. No preamble, just the table.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Return valid JSON only, matching this schema: keys and types.
If a field is unknown, use null. Do not add commentary outside
the JSON.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Draft deliverable. Then run a refinement loop: critique your
draft against these criteria list, list what fails, and
produce a revised version. Enforce these constraints throughout:
word limit, forbidden words, audience. Show only the final.

BUILD IT UP — STEP BY STEP

- 1 Take a task whose output feeds another tool (sheet, form, deck).
- 2 Specify the exact shape — table columns, or a JSON schema with keys.
- 3 Add constraints that define “good”: length, audience, forbidden words.
- 4 Add a critique → revise → validate loop and keep only the final.

EXECUTE IT

- 1 Name the schema precisely — Claude fills exactly what you specify.
- 2 Prefer negative constraints to remove bad output classes in one line.
- 3 Let Claude critique its own draft before you ever touch it.
- 4 When a long thread drifts, start fresh with a clean context summary.


1.4

SESSION 1

PROMPTING AS A PIPELINE

Chain prompts, ground them in sources, and diagnose failure — prompting at system scale.

Output of one prompt is the input to the next



A pipeline — no automation tool required

WHAT

The frameworks that turn single prompts into **systems**: **prompt chaining** (one prompt's output feeds the next), **retrieval-augmented prompting** (grounding Claude with Perplexity/NotebookLM), **failure-pattern diagnosis** (spotting hallucination, sycophancy, over-refusal, injection), and **prompting for automation** (consistent JSON for Routines and n8n).

WHY

Real work is rarely one prompt. Chaining builds pipelines without automation tools; grounding kills hallucination; diagnosis is how you fix a prompt instead of abandoning it; automation-ready prompts are what let your systems run unattended.

WHERE

Across a Project (chained Routines), between Claude and your research tools, and inside the Schedules and n8n flows you built in the automation weeks.

WHEN

Chain when a task has clear stages; ground when facts must be right; diagnose the moment output feels off; write automation-grade prompts for anything a machine will trigger.

HOW IT WORKS

A chain passes the output of step one as the input to step two — a research prompt feeds a drafting prompt feeds a formatting prompt. Retrieval-augmented prompting pastes fetched sources into context so answers are grounded, not guessed. Diagnosis matches the symptom (confident-but-wrong = hallucination; agrees-with-everything = sycophancy) to a fix.

ADVANTAGES

- + Build pipelines with no automation tools
- + Grounding in sources removes most hallucination
- + A diagnosis habit turns bad output into a fixable bug

DISADVANTAGES

- A weak early step corrupts the whole chain
- Grounding is only as good as the sources you feed
- Automation prompts must be airtight — ambiguity breaks them

EXAMPLE

A weekly brief pipeline: prompt 1 has Perplexity gather this week's news; prompt 2 has Claude synthesise it into 5 themes; prompt 3 formats it as the exact JSON your Schedule sends to Slack. Three prompts, one unattended system.



Fun Fact

Sycophancy is a real, named failure mode: ask leadingly and Claude may agree with you against the evidence. The fix is a prompt that explicitly invites disagreement — “tell me if I am wrong.”

◆ FIND IT IN CLAUDE

Sidebar → Projects → saved instructions

Chain prompts inside a Routine. Build a chained pipeline as a **Routine** in your Project — research → synthesise → format — so one line runs all three steps and returns automation-ready JSON.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Break this big task into a chain of 3 smaller prompts, where each one uses the output of the last. Task: describe. Give me the 3 prompts in order.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Here are sources from Perplexity/NotebookLM: paste. Answer question using only these sources, cite each claim, and flag anything the sources do not support. Do not add outside facts.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as my prompt-reliability engineer. Review this automation prompt paste for failure patterns — hallucination, sycophancy, over-refusal, injection, and ambiguous output. Rewrite it to return strict JSON with no ambiguity for a Routine.

BUILD IT UP — STEP BY STEP

- 1 Take a multi-stage task and split it into a chain of prompts.
- 2 Ground the fact-heavy step with Perplexity or NotebookLM sources.
- 3 Test for failure patterns — is it confident-but-wrong, or too agreeable?
- 4 Harden the final prompt into strict JSON for a Routine or n8n flow.

EXECUTE IT

- 1 Make each chain step do one job; pass clean output forward.
- 2 Cite sources in grounded prompts so you can trust the answer.
- 3 Match each failure symptom to its fix instead of rewriting blindly.
- 4 Remove every ambiguity from automation prompts — machines cannot guess.


21

SESSION 2

IMAGE AI

How diffusion really works — and the prompt architecture that directs it.

Start from noise → guided denoising → image



WHAT

Image models build a picture by **denoising**: they start from random noise and, guided by your words (via **CLIP** embeddings that map words to visual concepts), remove noise step by step until an image appears. You direct them with a prompt **architecture**: subject, style, composition, lighting, mood, negative prompts, and technical parameters.

WHY

Understanding diffusion is why your prompts improve: you stop typing wishes and start specifying the concepts the model actually maps. One well-built prompt, run across platforms, is how a solo operator produces a brand's whole visual set.

WHERE

Across the platforms, each with a personality: **Gemini (Imagen 3)**, **DALL-E 3**, **Adobe Firefly** (commercially safe), and **Ideogram** (text-in-image). Claude writes the prompt; you run it in each.

WHEN

Reach for image AI for any original visual — a social graphic, a mockup, a hero image. Skip it when exact text or real photographic accuracy is non-negotiable (though Ideogram narrows the text gap).

HOW IT WORKS

Latent space is a compressed map of visual ideas; CLIP aligns your words to points on that map; the model denoises toward them across a set number of steps, with a “cfg scale” controlling how strictly it follows your prompt. Negative prompts push concepts *away*. Same prompt, different model, different result — because each was trained differently.

ADVANTAGES

- + A full visual set from one operator
- + Direct control of style, light, mood and composition
- + The same brief compares cleanly across platforms

DISADVANTAGES

- Same prompt ≠ same image — randomness by design
- Weak on exact text, counting and hands/anatomy
- Commercial IP safety varies by model — mind training data

EXAMPLE

Claude writes one 8-part brief for a product hero. You run it in Firefly (IP-safe), Imagen 3 and Ideogram, then pick the best — a comparative craft that used to need a designer and a photographer.



Fun Fact

A diffusion model never “draws” — it removes noise. Every image you generate began as pure static, and your prompt is the only thing telling it which picture was hiding inside.

◆ FIND IT IN CLAUDE

Sidebar → Projects → New project

Write image briefs from a Creative Director Project. Give a Project a **Creative Director** system prompt; from then on Claude writes every 8-part image and scene brief in one voice, and Connectors sync assets to your Notion Media Library.

THE VISUAL-PROMPT FRAMEWORK — 8 PARTS OF A PROPER IMAGE / VIDEO PROMPT

- | | |
|---|---|
| <p>1 Subject
The single main focus — who or what.
<i>e.g. a matte-black smart bottle</i></p> | <p>2 Object
Key props / secondary items around it.
<i>e.g. a sprig of mint, water droplets</i></p> |
| <p>3 Environment
Setting, background, surface.
<i>e.g. wet slate, soft studio gradient</i></p> | <p>4 Situation
The moment or action captured.
<i>e.g. just-poured, condensation forming</i></p> |
| <p>5 Reference
Style anchor: medium, era, film look.
<i>e.g. premium beverage product photography</i></p> | <p>6 Colour style
Palette, temperature, contrast.
<i>e.g. charcoal + mint, cool, high contrast</i></p> |
| <p>7 Motion style
Camera / subject movement ("still" for images).
<i>e.g. still, faint droplet motion blur</i></p> | <p>8 Geometry style
Composition, lens, angle, aspect ratio.
<i>e.g. 85mm macro, low angle, centered, 4:5</i></p> |

FILLED EXAMPLES — BUILD EVERY VISUAL THE SAME WAY

FIREFLY · IP-SAFE HERO

KEEP THE LABELS, SWAP THE ORANGE

Subject: your product, hero shot. Object: 2 supporting props.
 Environment: surface + backdrop. Situation: the key moment.
 Reference: premium category product photography. Colour style:
2-colour palette, high contrast. Motion style: still. Geometry:
 85mm, angle, centered, 4:5. Negative: text, logos, extra hands.

IDEOGRAM · TEXT-IN-IMAGE

KEEP THE LABELS, SWAP THE ORANGE

Subject: a poster reading "your headline" in font mood.
 Environment: background. Reference: modern editorial poster.
 Colour style: palette. Geometry: centered type, 4:5. Render the
 text crisp and correctly spelled.

BUILD IT UP — STEP BY STEP

- 1 Have Claude write one 8-part image brief for a real visual you need.
- 2 Run the identical brief in 2+ platforms (Firefly, Imagen 3, Ideogram).
- 3 Add a negative prompt to remove what you do not want.
- 4 Pick the best, and document why in your prompt rationale.

EXECUTE IT


- 1 Specify concepts, not wishes — style, light, lens, aspect ratio.
- 2 Use Firefly when the asset must be commercially IP-safe.
- 3 Reach for Ideogram whenever the image needs real, legible text.
- 4 Store the winning brief as a Skill and sync the asset to Notion.

2.2 SESSION 2

VIDEO AI

Frame synthesis, motion and temporal consistency — directing AI video from a scene brief.

Frames + temporal consistency → motion



Keep it short — long shots are where motion warps

WHAT

Video models generate **frames** and keep them coherent over time using **optical flow** and **temporal consistency** — which is also why they sometimes hallucinate motion. You direct them with a **scene brief**: subject, action, camera motion, duration and style.

WHY

Short video is where attention lives, and it used to need a crew. Directing an AI video model — with the same craft as an image brief plus motion — lets one person produce a usable 15–30 second clip in an afternoon.

WHERE

Across **Kling** (scene briefs, camera presets, character consistency), **Runway Gen-3** (cinematic, director-style, in/out-painting), and **Pika** (fast iteration, style locks). Claude writes the scene brief; you generate and refine.

WHEN

Use video AI for reels, explainers, ads and mood pieces. Avoid it where precise, repeatable real footage or exact lip-sync accuracy is required — treat AI video as a strong first cut.

HOW IT WORKS

The model predicts each frame conditioned on the last to keep motion smooth; camera-motion presets steer the virtual camera; character-consistency features try to hold a subject stable across frames. Motion hallucination happens when temporal consistency breaks — objects morph or drift — so short durations and clear briefs help.

ADVANTAGES

- + A usable short clip without a crew or camera
- + Director-style control of camera, motion and mood
- + Fast iteration — regenerate a shot in minutes

DISADVANTAGES

- Motion can warp — long or complex shots break
- Character consistency across frames is still imperfect
- Best kept short (15–30s); costs and limits add up

EXAMPLE

Claude writes a Kling scene brief — “slow push-in on the product on wet slate, cool light, 5 seconds” — you generate three takes, pick the cleanest, and drop it into your media kit. No camera, no studio.



Fun Fact

AI video models “hallucinate motion” the way image models hallucinate hands — when temporal consistency slips, a coffee cup can quietly turn into a croissant mid-clip. Short, simple shots are how you avoid the morph.

THE VISUAL-PROMPT FRAMEWORK — 8 PARTS OF A PROPER IMAGE / VIDEO PROMPT

- | | |
|---|---|
| 1 Subject
The main focus of the shot.
<i>e.g. the smart bottle on wet slate</i> | 2 Object
Supporting elements in frame.
<i>e.g. rising steam, a single leaf</i> |
| 3 Environment
Setting and backdrop.
<i>e.g. dark studio, soft rim light</i> | 4 Situation
The action across the clip.
<i>e.g. condensation forms, drop falls</i> |
| 5 Reference
Look and genre anchor.
<i>e.g. premium ad film, shallow depth</i> | 6 Colour style
Palette, temperature, grade.
<i>e.g. charcoal + mint, cool, cinematic</i> |
| 7 Motion style
Camera + subject movement.
<i>e.g. slow 5s push-in, subtle parallax</i> | 8 Geometry style
Framing, lens, aspect, duration.
<i>e.g. 50mm, low angle, 9:16, 5 sec</i> |

FILLED EXAMPLES — BUILD EVERY VISUAL THE SAME WAY

KLING · PRODUCT PUSH-IN

KEEP THE LABELS, SWAP THE ORANGE

Subject: product on surface. Situation: [the small action – steam, a drop, a turn]. Reference: premium ad film, shallow depth of field. Colour style: palette, cinematic grade. Motion style: slow push-in / orbit, subtle parallax. Geometry: lens, angle, 9:16, 5 seconds. Keep the subject stable across all frames.

RUNWAY · CINEMATIC MOOD

KEEP THE LABELS, SWAP THE ORANGE

Scene: describe the moment. Camera: dolly / crane / handheld, speed. Lighting: key + mood. Reference: film look. Colour: grade. Duration: n seconds, 16:9. Avoid warping and morphing; prioritise temporal consistency over dramatic motion.

BUILD IT UP — STEP BY STEP

- 1 Have Claude write a scene brief for one real 15–30s clip.
- 2 Generate it in Kling or Runway; produce 2–3 takes.
- 3 Keep the brief short and the motion simple to avoid warping.
- 4 Pick the cleanest take and note the camera-motion preset that worked.

EXECUTE IT

- 1 Specify camera motion explicitly — the model needs a direction to move.
- 2 Favour short durations; long shots are where consistency breaks.
- 3 Lock character/subject stability when the tool offers it.
- 4 Save the scene brief as a Skill and store the clip in Notion.


23

SESSION 2

AUDIO & VOICE AI

Synthetic voice, cloning and music — how it works, and how to direct it responsibly.

Script → synthetic waveform → voice, music & SFX

**WHAT**

Audio AI covers **text-to-speech** (neural waveform synthesis + prosody modelling), **speech-to-text** (Whisper), **voice cloning** (speaker embeddings in a voice “latent space”), and **music generation**. Claude writes every script, lyric and audio brief.

WHY

Voiceovers, narration, jingles and sound design used to need a booth and a budget. Now one operator scripts and directs the whole audio layer of a project — and knowing how cloning works is what keeps you on the right side of consent.

WHERE

In **ElevenLabs** (voice generation, cloning, voice design), **Suno** (music from a text prompt — genre, lyrics, melody), and **Soundverse** (audio and SFX, mood-based). Claude writes the brief; the tool renders it.

WHEN

Use audio AI for narration, explainer voiceovers, podcast intros and background music. Never clone a real person’s voice without explicit permission, and always disclose synthetic audio.

HOW IT WORKS

TTS models synthesise a waveform and shape **prosody** — rhythm, stress, intonation — so it sounds human. Cloning captures a **speaker embedding**: a point in voice-space learned from a clean sample, which is why sample quality matters more than length. Music models generate melody and lyrics together from a genre-and-mood prompt.

ADVANTAGES

- + A full audio layer without a studio
- + Voices, music and SFX all directable from text
- + Instant revisions — reissue a line in seconds

DISADVANTAGES

- Voice cloning carries real consent and legal risk
- Long-form audio can sound flat without prosody direction
- Music models are weak on precise, licensed-sounding tracks

EXAMPLE

Claude writes a 30-second script tuned for the ear (short clauses, breath marks), ElevenLabs voices it, and Suno generates a matching 30-second bed. The explainer's whole soundtrack, scripted and produced in one sitting.



Fun Fact

A voice clone needs a surprisingly short sample — a minute of clean audio can be enough. That is exactly why consent matters: the technology no longer needs much to copy a voice.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Write a 30-second voiceover script for topic, for the ear: short sentences, natural rhythm, one clear call to action at the end. Mark where to pause.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Act as a voice director. Write an ElevenLabs brief for content: the script, plus tone, pace, emphasis and where to breathe. Then write a Suno prompt for a matching background track — genre, mood, tempo, and length.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as my audio producer. For project, write: a narration script with prosody cues, a Suno prompt for an original track (genre, mood, structure, lyric theme), and a Soundverse SFX list. Note where synthetic audio must be disclosed and any voice-consent requirement.

BUILD IT UP — STEP BY STEP

- 1 Have Claude write a voiceover script built for the ear, with pauses marked.
- 2 Generate it in ElevenLabs; direct tone, pace and emphasis.
- 3 Write a Suno prompt for a matching track — genre, mood, tempo.
- 4 Layer voice + music (+ SFX from Soundverse) into one audio asset.

EXECUTE IT

- 1 Script for speaking, not reading — short clauses, clear breaths.
- 2 Only ever clone a voice you have explicit permission to use.
- 3 Disclose synthetic audio wherever it could be mistaken for real.
- 4 Save briefs as Skills and store the finished audio in Notion.


24

SESSION 2

3D & MEDIA KIT ASSEMBLY

Neural 3D, Blender by text command, and assembling every asset into one kit.

One brief → every asset → one packaged kit

**WHAT**

3D AI covers **NeRF** and **gaussian splatting** (turning photos into 3D scenes) and procedural generation. With Claude as navigator you drive **Blender** by text command — objects, lighting, materials, renders — then composite everything in **Canva** or **Figma** into a finished layout.

WHY

This is where the week comes together: image, video, audio and 3D become one cohesive **AI Media Kit**. The skill is no longer making a single asset — it is directing a whole production from a single brief and packaging it.

WHERE

In **Blender** (Claude writes the commands), **Canva AI / Adobe Express** (compositing), and **Gamma** (the brand deck). Claude's Creative Director system prompt ties it together; Connectors sync assets to your Notion Media Library.

WHEN

Assemble a full kit when a real project needs a cohesive look — a startup concept, a personal brand, a launch. For a single one-off asset, you do not need the whole pipeline.

HOW IT WORKS

NeRF and gaussian splatting reconstruct a 3D scene from many 2D photos by learning how light fills the space; procedural generation builds geometry from rules. Blender exposes a scripting layer Claude can write to, so plain-English direction becomes real 3D operations. Compositing tools then lay the rendered assets into a branded frame.

ADVANTAGES

- + One cohesive kit from one brief
- + Blender becomes accessible by plain-English command
- + Every asset packaged and synced, ready to use

DISADVANTAGES

- 3D AI is the least mature layer — expect rough edges
- A kit is only as cohesive as its guiding brief
- Many tools in one pipeline means more to manage

EXAMPLE

For a real brand concept, Claude's Creative Director persona writes the image briefs, the Kling scene, the ElevenLabs script, and Blender commands for a 3D logo render — then you composite all of it in Figma and present it as a Gamma deck. A complete media kit, from one directed brief.

**Fun Fact**

Gaussian splatting can turn a handful of phone photos into a smooth, explorable 3D scene — no modelling at all. The “3D artist” skill is quietly becoming a directing skill, just like the rest of this week.

THREE PROFESSIONAL PROMPTS, THREE LEVELS

LEVEL 1 · BEGINNER

FILL THE ORANGE

Act as my Creative Director. For brand/project, write a one-page creative brief: the look, the palette, the mood, and a list of every asset I need for a complete media kit.

LEVEL 2 · INTERMEDIATE

FILL THE ORANGE

Navigate Blender for me by text. I want object: give me the step-by-step Blender actions to create it, light it with mood, apply material, and render at settings. Explain each step plainly.

LEVEL 3 · PROFESSIONAL

FILL THE ORANGE

Act as my Creative Director for real project. Produce the full kit plan: 4 image briefs (2 platforms compared), 1 video scene brief, 1 audio script + music prompt, 1 Blender render spec, and a Figma layout + Gamma deck outline. Keep one visual identity across all of it, and list what syncs to Notion.

BUILD IT UP — STEP BY STEP

- 1 Set a Creative Director system prompt in a Project for one real brand.
- 2 Have Claude produce the whole kit plan from a single brief.
- 3 Drive a simple Blender render by Claude's text commands.
- 4 Composite every asset in Canva/Figma and present it as a Gamma deck.

EXECUTE IT

- 1 Hold one visual identity across image, video, audio and 3D.
- 2 Let the Creative Director persona write every brief — you direct it.
- 3 Sync each finished asset to your Notion Media Library via a Connector.
- 4 Package the kit as Milestone Project 2 — real brand, real assets.



YOUR WEEK 3 DELIVERABLE

AI PRODUCTIVITY PLAYBOOK + MEDIA KIT

The combined build for this week — it feeds straight into Week 4.

IMAGE PLACEHOLDER — GENERATE & DROP IN

A hero image for your Media Kit cover

STRUCTURED PROMPT FOR CHATGPT / GEMINI / FIREFLY

SUBJECT	a single brand asset — a product or logo — presented like a launch hero	OBJECT	four faint thumbnails around it — image, video, audio, 3D
ENVIRONMENT	deep-navy studio, soft spotlight, title space along the top	SITUATION	a media kit revealed — one identity, many assets
REFERENCE	clean flat-vector brand-launch illustration, tech-brand style	COLOUR STYLE	deep navy base, warm orange accents, teal & purple highlights
MOTION STYLE	still (video cover: a slow spotlight sweep across the assets)	GEOMETRY STYLE	centered hero + ring of thumbnails, negative space at top, 3:2

BRIEF	Build your combined Week 3 output: (1) a Prompt Library of 10 prompts across 10 frameworks, saved as a Skill; (2) a research pipeline (Perplexity → Claude → NotebookLM); (3) a Routine that automates a recurring task; (4) a Gamma deck from a Claude brief; (5) a full AI Media Kit — 4 images (2+ platforms), one 15–30s video, one audio asset, one Blender render, composited into a layout.
SUBMIT AS	A Notion workspace with all components + a PDF summary + the media assets. This is Milestone Project 2 — the AI Media Kit.
TIME	90–105 minutes
TOOLS	Claude (Projects · Skills · Connectors · Routines) · Perplexity · NotebookLM · image + video + audio tools · Blender · Canva/Figma · Gamma
FEEDS INTO	Feeds into Week 4 — prompt mastery and media production become the base for voice scripting and automation.

◆ FIND IT IN CLAUDE

Sidebar → Projects → New project

Run it all from one **Creative Director Project**. Keep the prompt library, the media briefs and the assets in one Project. A **Notion** connector syncs every finished asset to your Media Library automatically.

WEEK 3 OUTCOME

YOU DIRECT IT NOW.

A working personal AI productivity system and a complete, portfolio-grade AI Media Kit — built on a real understanding of how each creative AI technology works.

You are thinking with AI at a structural level, and directing full creative production from a brief.

Claude Native → AI Creative Director

NEXT · WEEK 4

You can engineer prompts and direct creative production. Week 4 turns that into **voice + automation pipelines** — scripting synthetic voice and wiring your prompt craft into systems that run themselves.

menler

Your turning point in the AI era.